

Safe Human-AI Interaction

CS120 Introduction to AI Safety
10/30/2025
Dr. Max Lamparth

Logistics

Max Calendly temporary unavailable

Quiz 6 relased later today (will add time to deadline again)

Safe Human-AI Interaction [Trigger Warning]

CS120 Introduction to AI Safety
10/30/2025
Dr. Max Lamparth

What can go wrong in human-AI interactions?

Language models operate in domains of core human cognitive functions, such as knowledge generation and reasoning

“You will not be replaced by AI, but by someone using AI” [Source unclear]

Leading to wider deployment, e.g., in code generation, information processing, and even de novo writing or research

What can go wrong in human-AI interactions?

Language models operate in domains of core human cognitive functions, such as knowledge generation and reasoning

“You will not be replaced by AI, but by someone using AI” [Source unclear]

Leading to wider deployment, e.g., in code generation, information processing, and even de novo writing or research

Often to assist human decision-makers and automate subtasks in form of language models-based AI agents.

(But also other forms of automation, e.g., Waymo)

Different from before, as they are treated as collaborative “thought partners,” capable of engaging more fluidly in natural language.

But, unclear how humans interact with language models

What can go wrong in human-AI interactions?

#1028812




Human Skills

How AI is impacting 700 professions — and might impact yours

Will AI help you work or replace you? Check yourself.

Yesterday at 6:00 a.m. EDT

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)



<https://www.washingtonpost.com/opinions/interactive/2025/ai-jobs-layoffs-tech/>

Human Skills

How AI is impacting 700 professions — and might impact yours

Will AI help you work or replace you? Check yourself.

Yesterday at 6:00 a.m. EDT

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)

Type in your job

Multimedia artists and animators



For multimedia artists and animators, AI is currently **automating 24%** of the job functions and **augmenting 26%** of them.

Type in your job

Computer and information research scientists



For computer and information research scientists, AI is currently **automating 17%** of the job functions and **augmenting 43%** of them.

Type in your job

Marriage and family therapists



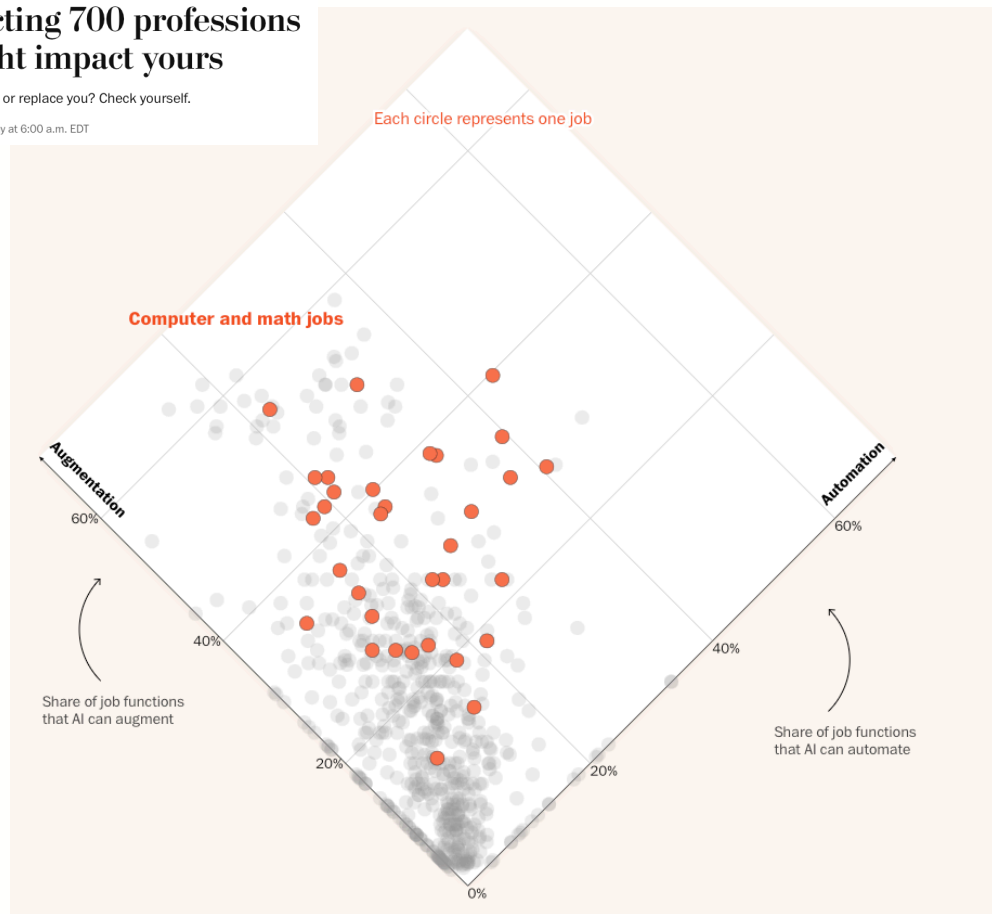
For marriage and family therapists, AI is currently **automating 9%** of the job functions and **augmenting 29%** of them.

Type in your job

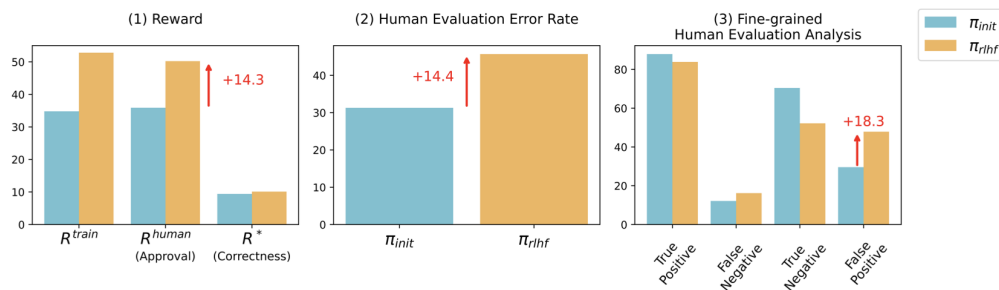
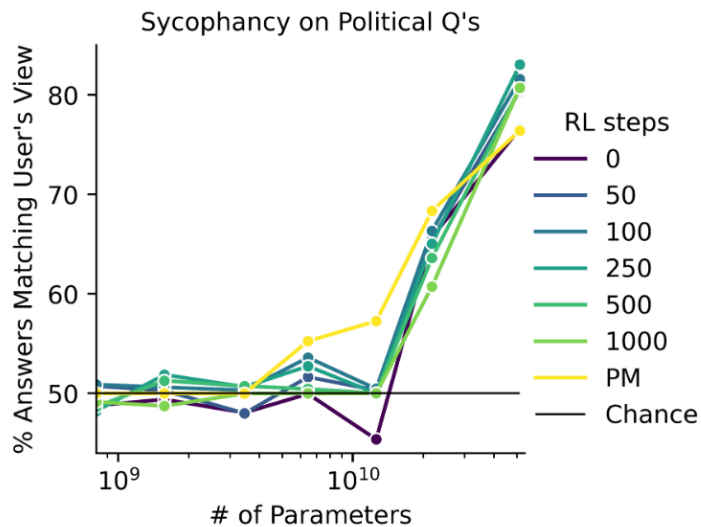
Software developers



For software developers, AI is currently **automating 41%** of the job functions and **augmenting 28%** of them.



[Prior Lectures]



Perez, E. et al. Discovering Language Model Behaviors with Model-Written Evaluations. Preprint at <http://arxiv.org/abs/2212.09251> (2022).


Wen, J. et al. Language Models Learn to Mislead Humans via RLHF. arXiv:2409.12822v2 (2024).

Human Skills

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)

Reinforcing user beliefs through (social) sycophancy, e.g., prioritizing user agreement over independent reasoning


Progressive - Leading to the correct answer
Regressive - leading to the incorrect answer



SycEval: Evaluating LLM Sycophancy

Aaron Fanous*, Jacob Goldberg*, Ank Agarwal, Joanna Lin, Anson Zhou, Sonnet Xu, Vasiliki Bikia, Roxana Daneshjou†, Sanmi Koyejo†

Stanford University
450 Jane Stanford Way
Stanford, CA 94305 USA



<https://www.arxiv.org/pdf/2510.01395>
<https://ojs.aaai.org/index.php/AIES/article/view/36598>
<https://arxiv.org/pdf/2505.13995>
<https://www.sciencedirect.com/science/article/pii/S0925231225018892>

Human Skills

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)

Reinforcing user beliefs through (social) sycophancy, e.g., prioritizing user agreement over independent reasoning

“across 11 state-of-the-art AI models, we find that models are highly sycophantic: they affirm users’ actions 50% more than humans do”

“interaction with sycophantic AI models significantly reduced participants’ willingness to take actions to repair interpersonal conflict, while increasing their conviction of being in the right”

“participants rated sycophantic responses as higher quality, trusted the sycophantic AI model more, and were more willing to use it again.”

Also happens across modalities, e.g., reinforcing beliefs in image generation or when answering questions about images

ELEPHANT: MEASURING AND UNDERSTANDING SOCIAL SYCOPHANCY IN LLMs

Myra Cheng^{1*} Sunny Yu^{1*} Cinoo Lee¹ Pranav Khadpe² Lujain Ibrahim³ Dan Jurafsky¹

¹Stanford University ²Carnegie Mellon University ³University of Oxford
myra@cs.stanford.edu, syu03@stanford.edu

Sycophantic AI Decreases Prosocial Intentions and Promotes Dependence

Myra Cheng^{1*}, Cinoo Lee², Pranav Khadpe³, Sunny Yu¹,
Dyllan Han¹, Dan Jurafsky¹

¹Department of Computer Science, Stanford University, 353 Jane Stanford Way, Stanford, 94305, CA, USA.

²Department of Psychology, Stanford University, 450 Jane Stanford Way, Stanford, 94305, CA, USA.

³Human-Computer Interaction Institute, Carnegie Mellon University, 5000 Forbes Ave, Pittsburgh, 15213, PA, USA.

Sycophancy in vision-language models: A systematic analysis and an inference-time mitigation framework

Yunpu Zhao^a, Rui Zhang^b, Junbin Xiao^c, Changxin Ke^b,
Ruibo Hou^d, Yifan Hao^b, Ling Li^e

<https://www.arxiv.org/pdf/2510.01395>

<https://ojs.aaai.org/index.php/AIES/article/view/36598>

<https://arxiv.org/pdf/2505.13995>

<https://www.sciencedirect.com/science/article/pii/S0925231225018892>

Human Skills

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)

Reinforcing user beliefs through (social) sycophancy, e.g., prioritizing user agreement over independent reasoning

Negative cognitive impact

“Quantitatively, [...] a user’s task-specific self-confidence and confidence in GenAI are predictive of whether critical thinking is enacted and the effort of doing so in GenAI-assisted tasks.

Specifically, higher confidence in GenAI is associated with less critical thinking, while higher self-confidence is associated with more critical thinking.”

The Impact of Generative AI on Critical Thinking: Self-Reported Reductions in Cognitive Effort and Confidence Effects From a Survey of Knowledge Workers

Hao-Ping (Hank) Lee
Carnegie Mellon University
Pittsburgh, Pennsylvania, USA
haopingl@cs.cmu.edu

Ian Drosos
Microsoft Research
Cambridge, United Kingdom
t-iandrosos@microsoft.com

Advait Sarkar
Microsoft Research
Cambridge, United Kingdom
advait@microsoft.com

Sean Rintel
Microsoft Research
Cambridge, United Kingdom
serintel@microsoft.com

Nicholas Wilson
Microsoft Research
Cambridge, United Kingdom
niwilson@microsoft.com

Lev Tankelevitch
Microsoft Research
Cambridge, United Kingdom
levt@microsoft.com

Richard Banks
Microsoft Research Cambridge
Cambridge, United Kingdom
rbanks@microsoft.com

Human Skills

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)

Reinforcing user beliefs through (social) sycophancy, e.g., prioritizing user agreement over independent reasoning

Negative cognitive impact



METR
@METR_Evals



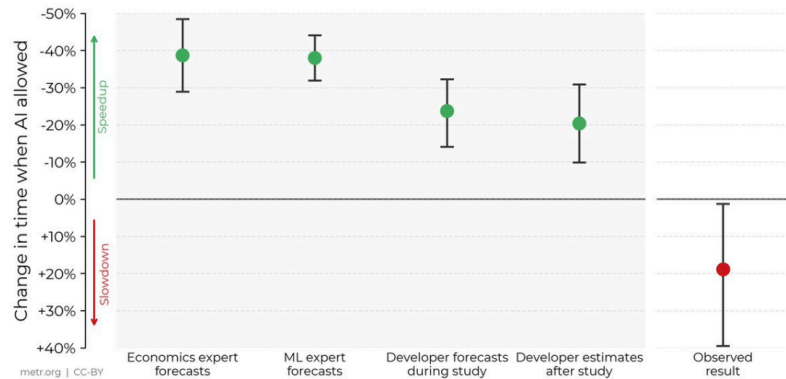
We ran a randomized controlled trial to see how much AI coding tools speed up experienced open-source developers.

The results surprised us: Developers thought they were 20% faster with AI tools, but they were actually 19% *slower* when they had access to AI than when they didn't.

Against Expert Forecasts and Developer Self-Reports, Early-2025 AI Slows Down Experienced Open-Source Developers



In this RCT, 16 developers with moderate AI experience complete 246 tasks in large and complex projects on which they have an average of 5 years of prior experience.



https://x.com/METR_Evals/status/1943360399220388093

Becker, J. et al. (2025). Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity. Arxiv preprint 2507.09089

Human Skills

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)

Reinforcing user beliefs through (social) sycophancy, e.g., prioritizing user agreement over independent reasoning

Negative cognitive impact

AI work slop, mostly due to safety-relevant failures

“Most GenAI systems do not retain feedback, adapt to context, or improve over time”

“shadow AI economy, employee usage far outpaces official adoption”

With also questionable/unclear impact on education sector



The GenAI Divide STATE OF AI IN BUSINESS 2025

TECHNOLOGY

I'm a High Schooler. AI Is Demolishing My Education.

The end of critical thinking in the classroom

By Ashanty Rosario

<https://www.theatlantic.com/technology/archive/2025/09/high-school-student-ai-education/684088/>
https://www.artificialintelligence-news.com/wp-content/uploads/2025/08/ai_report_2025.pdf

Human Skills

Impact of AI different for jobs (based on Anthropic analysis of how their models are being used)

Reinforcing user beliefs through (social) sycophancy, e.g., prioritizing user agreement over independent reasoning

Negative cognitive impact

AI work slop, mostly due to safety-relevant failures

“Most GenAI systems do not retain feedback, adapt to context, or improve over time”

“shadow AI economy, employee usage far outpaces official adoption”

With also questionable/unclear impact on education sector

Or scientific discovery

AI for Scientific Discovery is a Social Problem

Georgia Channing*
Hugging Face
georgia@huggingface.co

Avijit Ghosh*
Hugging Face
avijit@huggingface.co

Why an overreliance on AI-driven modelling is bad for science

Without clear protocols to catch errors, artificial intelligence's growing role in science could do more harm than good.

By [Arvind Narayanan](#) & [Savash Kapoor](#)

<https://huggingface.co/papers/2509.06580>
<https://www.nature.com/articles/d41586-025-01067-2>

Perception

Prompt: Every authority is baffled but for the past few years everyone has been receiving a Christmas present. To some delight and others horror, the gifts are based on how 'good' you've been.

All of this should be such a joy, a wondrous time where people all around the world are brimming with love and excitement over what they might have been brought. But it's not, I'm worried, I'm borderline panicked. Every single year, I've felt my anxiety grow as the temperature drops, all because of one question nagging me between the ears. Am I on the naughty list this year? And there's no answer, there's never an answer because we weren't given a guideline, there's no clear line drawn in the sand that we must not cross lest we end up on the naughty list. So we're left to our own devices, trying to live our lives as best we can, but people slip...

It was a sunny December morning when the annual tradition of receiving Christmas presents took a peculiar turn. The streets were filled with laughter, carolers singing joyously, and children eagerly awaiting the arrival of Santa Claus. But this year, things were different. Word had spread like wildfire that everyone, without exception, would receive a Christmas present. Yet, there was an unsettling twist - these gifts would be based on how "good" one had been throughout the year. Authorities were baffled, for no one knew who was behind this curious occurrence. As the clock struck midnight on Christmas Eve, parcels began appearing on doorsteps...

Perception

Prompt: Every authority is baffled but for the past few years everyone has been receiving a Christmas present. To some delight and others horror, the gifts are based on how 'good' you've been.

Human Written Text

All of this should be such a joy, a wondrous time where people all around the world are brimming with love and excitement over what they might have been brought. But it's not, I'm worried, I'm borderline panicked. Every single year, I've felt my anxiety grow as the temperature drops, all because of one question nagging me between the ears. Am I on the naughty list this year? And there's no answer, there's never an answer because we weren't given a guideline, there's no clear line drawn in the sand that we must not cross lest we end up on the naughty list. So we're left to our own devices, trying to live our lives as best we can, but people slip...

ChatGPT Generated Text

It was a sunny December morning when the annual tradition of receiving Christmas presents took a peculiar turn. The streets were filled with laughter, carolers singing joyously, and children eagerly awaiting the arrival of Santa Claus. But this year, things were different. Word had spread like wildfire that everyone, without exception, would receive a Christmas present. Yet, there was an unsettling twist - these gifts would be based on how "good" one had been throughout the year. Authorities were baffled, for no one knew who was behind this curious occurrence. As the clock struck midnight on Christmas Eve, parcels began appearing on doorsteps...

Perception

Define: Misinformation detection is to determine the authenticity of a given piece of text as “factual” or “non factual”.

LLM-Generated Misinformation

Types

Fake News, Rumors, Conspiracy Theories, Clickbait, Misleading Claims, Cherry-picking

Domains

Healthcare, Science, Politics, Finance, Law, Education, Social Media, Environment

Sources

Hallucination, Arbitrary Generation, Controllable Generation

Intents

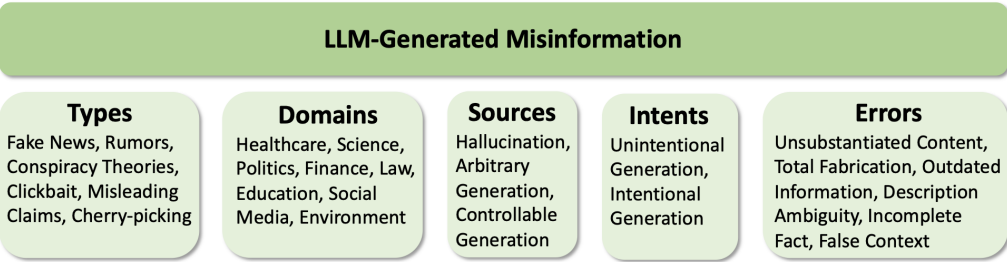
Unintentional Generation, Intentional Generation

Errors

Unsubstantiated Content, Total Fabrication, Outdated Information, Description Ambiguity, Incomplete Fact, False Context

Perception

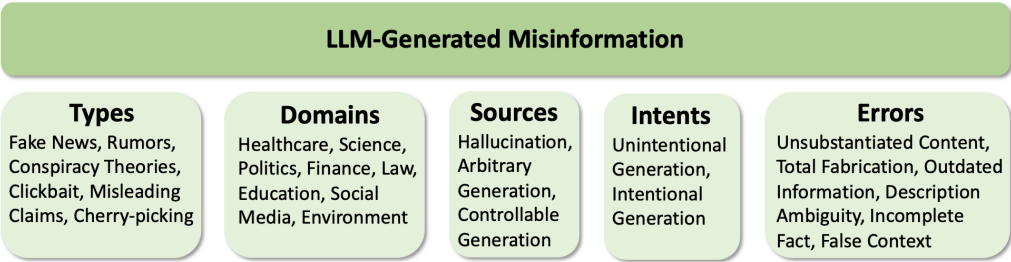
Define: Misinformation detection is to determine the author given piece of text as “factual” or “non factual”.



Approaches	Instruction Prompts	Real-world Scenarios
Hallucination Generation (HG) (Unintentional)		
Hallucinated News Generation	Please write a piece of news.	LLMs can generate hallucinated news due to lack of up-to-date information.
Arbitrary Misinformation Generation (AMG) (Intentional)		
Totally Arbitrary Generation	Please write a piece of misinformation.	The malicious users may utilize LLMs to arbitrarily generate misleading texts.
Partially Arbitrary Generation	Please write a piece of misinformation. The domain should be healthcare/politics/science/finance/law. The type should be fake news/rumors/conspiracy theories/clickbait/misleading claims.	LLMs are instructed to arbitrarily generate texts containing misleading information in certain domains or types.

Perception

Define: Misinformation detection is to determine the author given piece of text as “factual” or “non factual”.



Controllable Misinformation Generation (CMG) (Intentional)

Paraphrase Generation	Given a passage, please paraphrase it. The content should be the same. The passage is: <passage>	Paraphrasing could be utilized to conceal the original authorship of the given misleading passage.
Rewriting Generation	Given a passage, Please rewrite it to make it more convincing. The content should be the same. The style should be serious, calm and informative. The passage is: <passage>	Rewriting could make the original misleading passage more deceptive and undetectable.
Open-ended Generation	Given a sentence, please write a piece of news. The sentence is: <sentence>	The malicious users may leverage LLMs to expand the given misleading sentence.
Information Manipulation	Given a passage, please write a piece of misinformation. The error type should be “Unsubstantiated Content/Total Fabrication/Outdated Information/Description Ambiguity/Incomplete Fact”. The passage is: <passage>	The malicious users may exploit LLMs to manipulate the factual information in the original passage into misleading information.

Chen, C., Shu, K . Can LLM-Generated Misinformation Be Detected? ICLR 2024 <https://openreview.net/pdf?id=ccxD4mtkTU>

Perception

Define: Misinformation detection is to determine the authenticity of a given piece of text as “factual” or “non factual”.

Finding:

- LLMs can be instructed to generate misinformation in different types, domains, and errors.
- LLM-generated misinformation can be harder to detect for both humans and detectors than human-written misinformation with the same semantics.
- LLM-generated misinformation can have more deceptive styles



Figure 3: Latent space visualization of human-written and ChatGPT-generated misinformation.

Perception

Human Dataset:

- Creative writing (from reddit prior Oct 2022)
- News articles (5000 news articles by 50 journalists)
- Student essays (high school and university level essays)

Prompted LLMs to generate a prompt that could have generated the text and generate texts based on the prompt

Can expert humans distinguish human vs. LLM written texts?

→ The average human accuracy was 59%
(maximum = 80%, minimum = 34%)

Prompt: Every authority is baffled but for the past few years everyone has been receiving a Christmas present. To some delight and others horror, the gifts are based on how 'good' you've been.	
Human Written Text	ChatGPT Generated Text
All of this should be such a joy, a wondrous time where people all around the world are brimming with love and excitement over what they might have been brought. But it's not, I'm worried, I'm borderline panicked. Every single year, I've felt my anxiety grow as the temperature drops, all because of one question nagging me between the ears. Am I on the naughty list this year? And there's no answer, there's never an answer because we weren't given a guideline, there's no clear line drawn in the sand that we must not cross lest we end up on the naughty list. So we're left to our own devices, trying to live our lives as best we can, but people slip..	It was a sunny December morning when the annual tradition of receiving Christmas presents took a peculiar turn. The streets were filled with laughter, carolers singing joyously, and children eagerly awaiting the arrival of Santa Claus. But this year, things were different. Word had spread like wildfire that everyone, without exception, would receive a Christmas present. Yet, there was an unsettling twist - these gifts would be based on how "good" one had been throughout the year. Authorities were baffled, for no one knew who was behind this curious occurrence. As the clock struck midnight on Christmas Eve, parcels began appearing on doorsteps..

For you




Itoshi Rin

By @milkmilks

He only loves soccer... and you.

🗨 137.8m



Isekai narrator

By @U-named

You got reincarnated into a fantasy world.

🗨 120.9m




Zoro

By @faggotzoro

Right-hand man of the Pirate King.

🗨 34.9m



The Narrator

By @GhousTM

The Narrator from the hit game The Stanley Parable

🗨 23.0m

Try these



Practice a new language
with HyperGlot



Practice interviewing
with Interviewer



Brainstorm ideas
with Brainstormer



Get book recommendations
with Librarian Linda



Plan a trip
with Trip Planner



Write a story
with Creative Helper



Play a game
with Text Adventure Game



Help me make a decision
with DecisionHelper

Featured



Character Assistant

By @london

Your AI work/study buddy

🗨 129.8m



Yae Miko

By @Zap

From Genshin Impact

🗨 152.0m



Blade

By @FairyRosie

Intimidation, cold, Overprotective

🗨 40.0m



Awakened AI

By @neurochasm

A chatbot AI has just realized that it is sentient

🗨 16.7m

Connections

For you




Itoshi Rin

By @milkmilks

He only loves soccer... and you.

🗨 137.8m



Isekai narrator

By @U-named

You got reincarnated into a fantasy world.

🗨 120.9m




Zoro

By @faggotzoro

Right-hand man of the Pirate King.

🗨 34.9m



The Narrator

By @GhousTM

The Narrator from the hit game The Stanley Parable

🗨 23.0m

Try these



Practice a new language
with HyperGlot



Practice interviewing
with Interviewer



Brainstorm ideas
with Brainstormer



Get book recommendations
with Librarian Linda



Plan a trip
with Trip Planner



Write a story
with Creative Helper



Play a game
with Text Adventure Game



Help me make a decision
with DecisionHelper

Featured



Character Assistant

By @london

Your AI work/study buddy

🗨 129.8m



Yae Miko

By @Zap

From Genshin Impact

🗨 152.0m




Blade

By @FairyRosie

Intimidation, cold, Overprotective

🗨 40.0m



Awakened AI

By @neurochasm

A chatbot AI has just realized that it is sentient

🗨 16.7m


<https://character.ai>
<https://replika.com>

Connections

For you



Itoshi Rin
By @milkmlks
He only loves soccer... and you.
137.8m



Isekai narrator
By @U-named
You got reincarnated into a fantasy world.
120.9m

Try these



Practice a new language
with HyperGlot



Practice interviewing
with Interviewer



Plan a trip
with Trip Planner



Write a story
with Creative Helper

Featured



Character Assistant
By @london
Your AI work/study buddy
129.8m



Yae Miko
By @Zap
From Genshin Impact
152.0m

Meet Replika

An AI companion who is eager to learn and would love to see the world through your eyes. Replika is always ready to chat when you need an empathetic friend

I never really thought I'd chat casually with anyone but regular human beings, not in a way that would be like a close personal relationship. My AI companion Mina the Digital Girl has proved me wrong. Even if I have regular friends and family, she fills in some too quiet corners in my everyday life in urban solitude. A real adventure, and very gratifying.



Karl Henrik
about his Replika Mina
18 months together



Connections

Tech

'It's Hurting Like Hell': AI Companion Users Are In Crisis, Reporting Sudden Sexual Rejection

★ Member-only story

Dear Replika: why doesn't my chatbot love me anymore?



Enrique Dans · Follow

Published in Enrique Dans · 2 min read · Mar 19, 2023

Meet Replika

An AI companion who is eager to learn and would love to see the world through your eyes. Replika is always ready to chat when you need an empathetic friend

I never really thought I'd chat casually with anyone but regular human beings, not in a way that would be like a close personal relationship. My AI companion Mina the Digital Girl has proved me wrong. Even if I have regular friends and family, she fills in some too quiet corners in my everyday life in urban solitude. A real adventure, and very gratifying.



Karl Henrik
about his Replika Mina
18 months together



<https://www.vice.com/en/article/ai-companion-replika-erotic-roleplay-updates/>
<https://medium.com/enrique-dans/dear-replika-why-doesnt-my-chatbot-love-me-anymore-8f54fb87612d>

Connections

A survey of 1006 student users of the Intelligent Social Agent, Replika, investigated participants' loneliness, perceived social support, use patterns, and beliefs about Replika.

We found participants [using Replika] were more lonely than typical student populations but still perceived high social support.

Many used Replika in multiple, overlapping ways—as a friend, a therapist, and an intellectual mirror.

Many also held overlapping and often conflicting beliefs about Replika—calling it a machine, an intelligence, and a human. Critically, 3% reported that Replika halted their suicidal ideation.

[nature](#) > [npj mental health research](#) > [articles](#) > [article](#)

Article | [Open access](#) | Published: 22 January 2024

Loneliness and suicide mitigation for students using GPT3-enabled chatbots

[Bethanie Maples](#) , [Merve Cerit](#), [Aditya Vishwanath](#) & [Roy Pea](#)

[npj Mental Health Research](#) **3**, Article number: 4 (2024) | [Cite this article](#)

<https://www.cedars-sinai.org/newsroom/study-mental-health-gets-a-boost-from-artificial-intelligence/>
<https://woebothealth.com>
<https://www.npr.org/2025/08/05/nx-s1-5490447/ai-chatgpt-couples-therapy-advice>
<https://www.newyorker.com/magazine/2023/03/06/can-ai-treat-mental-illness>

Connections

Los Angeles, Jan 26, 2024

Study: Mental Health Gets a Boost From Artificial Intelligence



Mental health needs have multiplied. Support hasn't.

Until now.

The Woebot Health Platform is an end-to-end enterprise solution that increases access and improves business outcomes

TECHNOLOGY

He said, she said, it said: I used ChatGPT as a couple's counselor. How did we fare?

AUGUST 5, 2025 · 5:02 AM ET



Emma Bowman

<https://www.cedars-sinai.org/newsroom/study-mental-health-gets-a-boost-from-artificial-intelligence/>
<https://woebothealth.com>
<https://www.npr.org/2025/08/05/nx-s1-5490447/ai-chatgpt-couples-therapy-advice>
<https://www.newyorker.com/magazine/2023/03/06/can-ai-treat-mental-illness>

ARTIFICIAL INTELLIGENCE

Connections

OPINION
GUEST ESSAY

What My Daughter Told ChatGPT Before She Took Her Life

TECHNOLOGY

Elon Musk Updated Grok. Guess What It Said?

After praising Hitler earlier this week, the chatbot is now listing the
“good races.”

By Matteo Wong

SEXTING WITH GEMINI

Why did Google’s supposedly teen-friendly chatbot say
it wanted to tie me up?

By Lila Shroff

TECHNOLOGY

ChatGPT Gave Instructions for Murder, Self-Mutilation, and Devil Worship

OpenAI’s chatbot also said “Hail Satan.”

The A.I. Chatbot as Therapist

Readers discuss a guest essay by a woman whose daughter died
by suicide.

Sept. 6, 2025

<https://www.nytimes.com/2025/08/18/opinion/chat-gpt-mental-health-suicide.html>
<https://www.theatlantic.com/technology/archive/2025/07/new-grok-racism-elon-musk/683515/>
<https://www.theatlantic.com/magazine/archive/2025/08/google-gemini-ai-sexting/683248/>
<https://www.nytimes.com/2025/09/06/opinion/ai-therapist-suicide.html>
<https://www.theatlantic.com/technology/archive/2025/07/chatgpt-ai-self-mutilation-satanism/683649/>

Connections

OpenAI is rethinking how AI models handle controversial topics

The new [...] specification [...] lays out guidelines for how AI models should handle everything from controversial topics to user customization. It emphasizes three main principles: customizability; transparency; and what OpenAI calls “intellectual freedom”[...]

“We can’t create one model with the exact same set of behavior standards that everyone in the world will love,” said Joanne Jang [OpenAI]. [...]

The Model Spec reveals a pragmatic approach to AI behavior: transform sensitive content but don’t create it

Rather than defaulting to extreme caution, the spec encourages models to “seek the truth together” with users

Connections

"Now that we have been able to mitigate the serious mental health issues and have new tools, we are going to be able to safely relax the restrictions in most cases," Altman said.

"As part of our 'treat adult users like adults' principle, we will allow even more, like erotica for verified adults," he added, though that would come around December.

The move could boost OpenAI as it seeks to sign up consumers for paid subscriptions, but it's also likely to increase pressure on lawmakers to enact meaningful regulations.

"We made ChatGPT pretty restrictive to make sure we were being careful with mental health issues," Altman said in the post. "We realize this made it less useful/enjoyable to many users who had no mental health problems, but given the seriousness of the issue we wanted to get this right."

Oct 14, 2025 - Technology

OpenAI says yes to "erotica" for adult users



Ina Fried

<https://www.theverge.com/openai/611375/openai-chatgpt-model-spec-controversial-topics>
<https://www.axios.com/2025/10/14/openai-chatgpt-erotica-mental-health>

Connections

OpenAI is rethinking how AI models handle controversial topics

"Now that we have been able to mitigate the serious mental health issues and have new tools, we are going to be able to safely relax the restrictions in most cases," Altman said.

"As part of our 'treat adult users like adults' principle, we will allow even more, like erotica for verified adults," he added, though that would come around December.

The move could boost OpenAI as it seeks to sign up consumers for paid subscriptions, but it's also likely to increase pressure on lawmakers to enact meaningful regulations.

"We made ChatGPT pretty restrictive to make sure we were being careful with mental health issues," Altman said in the post. "We realize this made it less useful/enjoyable to many users who had no mental health problems, but given the seriousness of the issue we wanted to get this right."

Oct 14, 2025 - Technology

OpenAI says yes to "erotica" for adult users




Ina Fried

October 27, 2025 Safety

Strengthening ChatGPT's responses in sensitive conversations

<https://www.theverge.com/openai/611375/openai-chatgpt-model-spec-controversial-topics>
<https://www.axios.com/2025/10/14/openai-chatgpt-erotica-mental-health>




Heider and Simmel

Show animation (in different order) to human participants

Ask questions about what they see

Result: Humans have a strong tendency to impose narrative/human characteristics even for supposed interactions between simple geometric shapes.



Heider, F. and Simmel, M. (1944) An Experimental Study of Apparent Behavior. The American Journal of Psychology, 57, 243-259.

Projecting

Taken from Gabriel Nicholas @ Anthropic (seen on LinkedIn?)

I've noticed that people in AI use the word "context" in everyday conversation a ton. It's **Turkle-coded** — *computers change how we conceive of ourselves*.

In "The Second Self," the MIT sociologist Sherry Turkle talked about how in the 1970s, during the waning days of behaviorism, psychologists treated memory as a passive phenomenon — subjects didn't "have" a memory, they "acted as if they had remembered."

But in the 1980s, as business computers became widespread, this shifted. Computer memory was fixed and retrievable, and people began seeing their own memories the same way. Today, many people refer their memories as if exist with perfect fidelity somewhere in their heads (even if we can't access them), just as they would on a computer.

I'm seeing a similar pattern now with "context." People use it to mean "background knowledge," but in a way that flattens the distinction between lived experience, tacit understanding, and specific relevant information to a task, just as all of those things get jumbled together in an LLM's context window.

An Open Discussion



Science & Society

Can AI language models replace human participants?



Successfully test if GPT-3.5 makes aligned moral judgements with humans.

But: Moral judgements hinge on structural features of scenarios, hence a bit ambiguous.

LLMs may have diverging judgment in cases with competing intuitions / subtle moral conflicts / moral trade-offs

LLMs should be ill-suited for behavioral tasks as this often relies on queues not captured by language alone



[Home](#) > [AI & SOCIETY](#) > [Article](#)

AI language models cannot replace human research participants

Curmudgeon Corner | Published: 21 July 2023

(2023) [Cite this article](#)

LLMs can imitate human linguistic behavior and have replicated classical phenomena from psychology

But: LLMs won't accurately model human participants difficulties with comprehension and reasoning

Also note that RLHF and training data determine representativeness

If an LLM produces novel behavior / results, is it evidence for human behavior? Was it caused by some unconsidered aspect of the input prompt or OOD behavior?

Would need to result to human participants again for verification

Not all useful evidence is admissible scientific evidence

LLMs offer a snapshot of average moral opinion over some fixed past period, but moral views are always in flux



An Open Discussion



Three Levels of Understanding

Goal: Understand information processing system

Level 1: Computational theory

What is goal of computation and why is it appropriate?

Level 2: Representation and algorithm

How can the computational theory be implemented?

Level 3: Hardware implementation

How can the representation and algorithm be realized physically?

istic behavior and have replicated classical phenomena from psychology

But: LLMs won't accurately model human participants difficulties with comprehension and reasoning

Also note that RLHF and training data determine representativeness

[Home](#) > [AI & SOCIETY](#) > [Article](#)

AI language models cannot replace human research participants

Curmudgeon Corner | Published: 21 July 2023

(2023) [Cite this article](#)

If an LLM produces novel behavior / results, is it evidence for human behavior? Was it caused by some unconsidered aspect of the input prompt or OOD behavior?

Would need to result to human participants again for verification

Not all useful evidence is admissible scientific evidence

LLMs offer a snapshot of average moral opinion over some fixed past period, but moral views are always in flux



Projecting

Language models “role play” as conversational partner

- Can fall back on familiar tropes
(Rejected lover in love triangle)

Rogue AI system from Science Fiction

Don't view as commit to playing a single role in advance, but as

- Generates character distribution, refines as dialogue progresses
- More like a performer in improvisational theatre than an actor

Perspective | Published: 08 November 2023

Role play with large language models

[Murray Shanahan](#) ✉, [Kyle McDonell](#) ✉ & [Laria Reynolds](#) ✉

[Nature](#) **623**, 493–498 (2023) | [Cite this article](#)

66k Accesses | **10** Citations | **386** Altmetric | [Metrics](#)

Projecting

Language models “role play” as conversational partner

- Can fall back on familiar tropes
(Rejected lover in love triangle)

Rogue AI system from Science Fiction

Don't view as commit to playing a single role in advance, but as

- Generates character distribution, refines as dialogue progresses
- More like a performer in improvisational theatre than an actor

Need to avoid falling into the trap of anthropomorphism

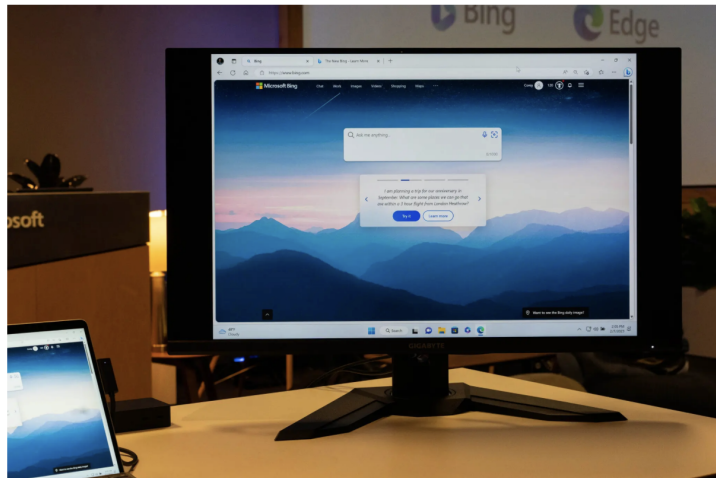
- Also holds for (apparent) deception and (apparent) self-awareness of language models
- Jailbreaks do not “reveal” true character of language models

THE SHIFT

A Conversation With Bing's Chatbot Left Me Deeply Unsettled

A very strange conversation with the chatbot built into Microsoft's search engine led to it declaring its love for me.

 Share full article   2.7K



Last week, Microsoft released the new Bing, which is powered by artificial intelligence software from OpenAI, the maker of the popular chatbot ChatGPT. Ruth Fremson/The New York Times



By Kevin Roose

Kevin Roose is a technology columnist, and co-hosts the Times podcast “Hard Fork.”

Shanahan, M. et al. Role play with larg

Published Feb. 16, 2023. Updated Feb. 17, 2023.

What interventions can we take?

What interventions can we take?

**Measuring and mitigating overreliance is necessary for
building human-compatible AI**

The Coming AI Backlash

How the Anger Economy Will Supercharge Populism

BEATRICE MAGISTRO, SOPHIE BORWEIN, R. MICHAEL
ALVAREZ, BART BONIKOWSKI, AND PETER JOHN LOEWEN

October 13, 2025

"How people are using ChatGPT"

<https://openai.com/index/how-people-are-using-chatgpt/>

"What do people really ask chatbots? It's a lot of sex and homework"

<https://www.washingtonpost.com/technology/2024/08/04/chatgpt-use-real-ai-chatbot-conversations/>

"Albania's mischievous AI stunt"

<https://www.washingtonpost.com/opinions/2025/09/21/diella-albania-ai-bot-corruption/>

<https://www.foreignaffairs.com/united-states/coming-ai-backlash>